

Pipeline

March 7, 2017

0.1 Advanced Lane Finding Project

```
In [1]: import cv2
import os
import numpy as np
import matplotlib.pyplot as plt
from scipy.misc import imread as imread
from scipy.misc import imsave as imsave
from scipy.misc import imresize as imresize
from IPython.display import HTML
from moviepy.editor import VideoFileClip
from moviepy.editor import ImageSequenceClip
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: %matplotlib inline
```

0.2 Camera Calibration

```
In [5]: # Load chessboard images for camera calibration
```

```
folder = 'camera_cal/'
image_files = os.listdir(folder)
chessboard_images = []

for i in range(len(image_files)):
    file_name = os.path.join(folder, ('calibration' + str(i+1) + '.jpg'))
    image_file = imread(file_name, False, 'RGB')
    image_file = imresize(image_file, (720, 1280), interp='bilinear')
    chessboard_images.append(image_file)
```

```
In [6]: # Compute the camera calibration matrix and distortion coefficients
# given a set of chessboard images
```

```
def compute_cal_dist(img):
    # Prepare object points, like (0,0,0), (1,0,0), (2,0,0) ..., (6,5,0)
    objp = np.zeros((6*8,3), np.float32)
    objp[:, :2] = np.mgrid[0:8, 0:6].T.reshape(-1,2)
```

```

    # Convert to grayscale
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

    # Find the chessboard corners
    ret, corners = cv2.findChessboardCorners(gray, (8,6), None)

    objpoints = objp
    imgpoints = corners

    return ret, objpoints, imgpoints

In [7]: # Arrays to store object points and image points from all the images.
    objpoints = [] # 3d points in real world space
    imgpoints = [] # 2d points in image plane.

    for i in range(20):
        img = chessboard_images[i]
        ret, objp, imp = compute_cal_dist(img)

        # If found, add object points, image points
        if ret == True:
            objpoints.append(objp)
            imgpoints.append(imp)

In [9]: # Do camera calibration given object points and image points
    img = cv2.imread('test_image.jpg')
    img_size = (img.shape[1], img.shape[0])

    # Do camera calibration given object points and image points
    ret, mtx, dist, rvecs, tvecs = cv2.calibrateCamera(objpoints, imgpoints, img_size, None, None)

In [10]: def calibrate_image(image, mtx, dist):
    '''
        Applies distortion correction
        '''

        # Apply distortion correction to raw images.
        undst = cv2.undistort(image, mtx, dist, None, mtx)

        return undst

```

0.3 Pipeline Functions

```

In [11]: def mag_thresh(img, sobel_kernel=3, mag_thresh=(0, 255)):
    '''
        To create an edge image
        '''

        # Convert to grayscale

```

```

gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)

# Take both Sobel x and y gradients
sobelx = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize=sobel_kernel)
sobely = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize=sobel_kernel)

# Calculate the gradient magnitude
gradmag = np.sqrt(sobelx**2 + sobely**2)

# Rescale to 8 bit
scale_factor = np.max(gradmag)/255
gradmag = (gradmag/scale_factor).astype(np.uint8)

# Create a binary image of ones where threshold is met, zeros otherwise
binary_output = np.zeros_like(gradmag)
binary_output[(gradmag >= mag_thresh[0]) & (gradmag <= mag_thresh[1])] = 1

# Return the binary image
return binary_output

```

```

In [12]: def grad_threshold(img, sobel_kernel=3, thresh=(0, np.pi/2)):
    '''
    To create a gradient image
    '''
    # Grayscale
    gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)

    # Calculate the x and y gradients
    sobelx = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize=sobel_kernel)
    sobely = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize=sobel_kernel)

    # Take the absolute value of the gradient direction,
    # apply a threshold, and create a binary image result
    absgraddir = np.arctan2(np.absolute(sobely), np.absolute(sobelx))
    binary_output = np.zeros_like(absgraddir)
    binary_output[(absgraddir >= thresh[0]) & (absgraddir <= thresh[1])] = 1

    # Return the binary image
    return binary_output

```

```

In [13]: def color(img, thresh=(0, 255)):
    '''
    To create a color map image
    '''
    # Convert to HLS color space
    hls = cv2.cvtColor(img, cv2.COLOR_RGB2HLS)

    # Select the S-channel

```

```

S = hls[:, :, 2]

# Take the HLS color space, apply a threshold, and create a binary image
binary = np.zeros_like(S)
binary[(S > thresh[0]) & (S <= thresh[1])] = 1

return binary

```

```

In [14]: def region_of_interest(img, vertices):
    '''
    Applies an image mask.
    Only keeps the region of the image defined by the polygon
    formed from `vertices`. The rest of the image is set to black.
    '''
    # Defining a blank mask to start with
    mask = np.zeros_like(img)

    # Defining a 3 channel or 1 channel color to fill the mask with depending on image
    if len(img.shape) > 2:
        channel_count = img.shape[2] # i.e. 3 or 4 depending on your image
        ignore_mask_color = (255,) * channel_count
    else:
        ignore_mask_color = 255

    # Filling pixels inside the polygon defined by "vertices" with the fill color
    cv2.fillPoly(mask, vertices, ignore_mask_color)

    # Returning the image only where mask pixels are nonzero
    masked_image = cv2.bitwise_and(img, mask)

    return masked_image

```

```

In [15]: def corners_unwarp(img, vx1, vx2, vy1, vy2):
    '''
    Perspective Transform
    '''
    # Input should be the calibrated Image
    undist = img

    # Grab the image shape
    img_size = (img.shape[1], img.shape[0])

    # For source points:
    src = np.float32([vx1, vx2, vy1, vy2])

    # For destination points:
    dst = np.float32([vx1, [vx1[0], 0], [vy2[0], 0], vy2])

```

```

        # Given src and dst points, calculate the perspective transform matrix
        M = cv2.getPerspectiveTransform(src, dst)

        # Compute the inverse perspective transform:
        Minv = cv2.getPerspectiveTransform(dst, src)

        # Warp the image using OpenCV warpPerspective()
        warped = cv2.warpPerspective(undist, M, img_size)

        # Return the resulting image and matrix
        return warped, M, Minv

In [16]: def binary_map_norm(img):
        '''
        To normalize the binary image map
        '''

        # Find the min value in the array
        img = np.array(img)
        min_value = np.min(img)

        # Loop through each element of the image to replace non-zero elements
        for i in range(img.shape[0]):
            for j in range(img.shape[1]):
                if img[i][j] >= min_value and img[i][j] != 0:
                    img[i][j] = 255.0

        img = img/255.0

        return img

In [17]: def curvature(top_down, left_lane_inds, right_lane_inds, left_fit, right_fit):
        '''
        Find the lane curvature
        '''

        # Identify the x and y positions of all nonzero pixels in the image
        nonzero = top_down.nonzero()
        nonzeroy = np.array(nonzero[0])
        nonzerox = np.array(nonzero[1])

        lefty = nonzeroy[left_lane_inds]
        righty = nonzeroy[right_lane_inds]

        y_eval1 = np.max(lefty)
        y_eval2 = np.max(righty)

        left_curverad = ((1 + (2*left_fit[0]*y_eval1 + left_fit[1])**2)**1.5)
        right_curverad = ((1 + (2*right_fit[0]*y_eval2 + right_fit[1])**2)**1.5)

        return left_curverad, right_curverad

```

```

In [18]: def first_lane(image, vertices):
    '''
    To find the lane in the first frame of video sequence.
    '''
    # Apply binary map to the next images.
    b_map = mag_thresh(image, sobel_kernel=3, mag_thresh=(40, 150))

    # Gradient Orientation
    grad = grad_threshold(image, sobel_kernel=15, thresh=(0.7, 1.3))

    # Color
    color_map = color(image, thresh=(90, 255))

    # Combine
    combined = np.zeros_like(grad)
    combined[((b_map == 1)) | ((color_map == 1) & (grad == 1))] = 1

    masked_edge_image = region_of_interest(combined, vertices)
    masked_image = region_of_interest(image, vertices)

    top_down, M, Minv = corners_unwarp(masked_edge_image, vertices[0,0,:],
    top_down = binary_map_norm(top_down)

    # Compute the histogram
    histogram = np.sum(top_down[top_down.shape[0]/2:,:], axis=0)

    # Find the peak of the left and right halves of the histogram
    # These will be the starting point for the left and right lines
    midpoint = np.int(histogram.shape[0]/2)
    leftx_base = np.argmax(histogram[:midpoint])
    rightx_base = np.argmax(histogram[midpoint:]) + midpoint

    # Choose the number of sliding windows
    nwindows = 9

    # Set height of windows
    window_height = np.int(top_down.shape[0]/nwindows)

    # Identify the x and y positions of all nonzero pixels in the image
    nonzero = top_down.nonzero()
    nonzeroy = np.array(nonzero[0])
    nonzerox = np.array(nonzero[1])

    # Current positions to be updated for each window
    leftx_current = leftx_base
    rightx_current = rightx_base

    # Set the width of the windows +/- margin

```

```

margin = 100

# Set minimum number of pixels found to recenter window
minpix = 50

# Create empty lists to receive left and right lane pixel indices
left_lane_inds = []
right_lane_inds = []

# Step through the windows one by one
for window in range(nwindows):
    # Identify window boundaries in x and y (and right and left)
    win_y_low = top_down.shape[0] - (window+1)*window_height
    win_y_high = top_down.shape[0] - window*window_height
    win_xleft_low = leftx_current - margin
    win_xleft_high = leftx_current + margin
    win_xright_low = rightx_current - margin
    win_xright_high = rightx_current + margin

    # Identify the nonzero pixels in x and y within the window
    good_left_inds = ((nonzeroy >= win_y_low) & (nonzeroy < win_y_high))
    good_right_inds = ((nonzeroy >= win_y_low) & (nonzeroy < win_y_high))

    # Append these indices to the lists
    left_lane_inds.append(good_left_inds)
    right_lane_inds.append(good_right_inds)

    # If you found > minpix pixels, recenter next window on their mean
    if len(good_left_inds) > minpix:
        leftx_current = np.int(np.mean(nonzerox[good_left_inds]))
    if len(good_right_inds) > minpix:
        rightx_current = np.int(np.mean(nonzerox[good_right_inds]))

# Concatenate the arrays of indices
left_lane_inds = np.concatenate(left_lane_inds)
right_lane_inds = np.concatenate(right_lane_inds)

# Extract left and right line pixel positions
leftx = nonzerox[left_lane_inds]
lefty = nonzeroy[left_lane_inds]
rightx = nonzerox[right_lane_inds]
righty = nonzeroy[right_lane_inds]

# Fit a second order polynomial to each
left_fit = np.polyfit(lefty, leftx, 2)
right_fit = np.polyfit(righty, rightx, 2)

# Find the curve:

```

```

left_curverad, right_curverad = curvature(top_down, left_lane_inds, ri

# Generate x and y values for plotting
ploty = np.linspace(0, top_down.shape[0]-1, top_down.shape[0] )
left_fitx = left_fit[0]*ploty**2 + left_fit[1]*ploty + left_fit[2]
right_fitx = right_fit[0]*ploty**2 + right_fit[1]*ploty + right_fit[2]

# Create an image to draw the lines on
font = cv2.FONT_HERSHEY_SIMPLEX
warp_zero = np.zeros_like(masked_edge_image).astype(np.uint8)
color_warp = np.dstack((warp_zero, warp_zero, warp_zero))

# Recast the x and y points into usable format for cv2.fillPoly()
pts_left = np.array([np.transpose(np.vstack([left_fitx, ploty]))])
pts_right = np.array([np.flipud(np.transpose(np.vstack([right_fitx, p
pts = np.hstack((pts_left, pts_right))

# Draw the lane onto the warped blank image
cv2.fillPoly(color_warp, np.int_([pts]), (0,255, 0))
cv2.putText(image, 'Radius of Curvature: ' + str(left_curverad) + 'm',

# Warp the blank back to original image space using inverse perspective
newwarp = cv2.warpPerspective(color_warp, Minv, (image.shape[1], image

# Combine the result with the original image
result = cv2.addWeighted(image, 1, newwarp, 0.3, 0)

return result, left_fit, right_fit

```

```

In [19]: def second_to_end_lane(image, vertices, left_fit, right_fit):
'''
To find the lane from second frame to end
'''

# Apply binary map to the next images.
b_map = mag_thresh(image, sobel_kernel=3, mag_thresh=(40, 150))

# Gradient Orientation
grad = grad_threshold(image, sobel_kernel=15, thresh=(0.7, 1.3))

# Color
color_map = color(image, thresh=(90, 255))

# Combine
combined = np.zeros_like(grad)
combined[((b_map == 1)) | ((color_map == 1) & (grad == 1))] = 1

masked_edge_image = region_of_interest(combined, vertices)

```



```

masked_image = region_of_interest(image, vertices)

top_down, M, Minv = corners_unwarp(masked_edge_image, vertices[0,0,:],
top_down = binary_map_norm(top_down)

nonzero = top_down.nonzero()
nonzero_y = np.array(nonzero[0])
nonzero_x = np.array(nonzero[1])
margin = 100
left_lane_inds = ((nonzero_x > (left_fit[0]*(nonzero_y**2) + left_fit[1]
right_lane_inds = ((nonzero_x > (right_fit[0]*(nonzero_y**2) + right_fit[1]

# Again, extract left and right line pixel positions
leftx = nonzero_x[left_lane_inds]
lefty = nonzero_y[left_lane_inds]
rightx = nonzero_x[right_lane_inds]
righty = nonzero_y[right_lane_inds]

# Fit a second order polynomial to each
left_fit = np.polyfit(lefty, leftx, 2)
right_fit = np.polyfit(righty, rightx, 2)

# Generate x and y values for plotting
ploty = np.linspace(0, top_down.shape[0]-1, top_down.shape[0] )
left_fitx = left_fit[0]*ploty**2 + left_fit[1]*ploty + left_fit[2]
right_fitx = right_fit[0]*ploty**2 + right_fit[1]*ploty + right_fit[2]

# Generate a polygon to illustrate the search window area
# And recast the x and y points into usable format for cv2.fillPoly()
left_line_window1 = np.array([np.transpose(np.vstack([left_fitx-margin,
left_line_window2 = np.array([np.flipud(np.transpose(np.vstack([left_
left_line_pts = np.hstack((left_line_window1, left_line_window2))
right_line_window1 = np.array([np.transpose(np.vstack([right_fitx-margin,
right_line_window2 = np.array([np.flipud(np.transpose(np.vstack([right_
right_line_pts = np.hstack((right_line_window1, right_line_window2))

# Find the curve:
left_curverad, right_curverad = curvature(top_down, left_lane_inds, ri

# Create an image to draw the lines on
font = cv2.FONT_HERSHEY_SIMPLEX
warp_zero = np.zeros_like(masked_edge_image).astype(np.uint8)
color_warp = np.dstack((warp_zero, warp_zero, warp_zero))

# Recast the x and y points into usable format for cv2.fillPoly()
pts_left = np.array([np.transpose(np.vstack([left_fitx, ploty]))])
pts_right = np.array([np.flipud(np.transpose(np.vstack([right_fitx, pl
pts = np.hstack((pts_left, pts_right))

```

```

# Draw the lane onto the warped blank image
cv2.fillPoly(color_warp, np.int_([pts]), (0,255, 0))
cv2.putText(image, 'Radius of Curvature: ' + str(left_curverad) + 'm',

# Warp the blank back to original image space using inverse perspective
newwarp = cv2.warpPerspective(color_warp, Minv, (image.shape[1], image

# Combine the result with the original image
result = cv2.addWeighted(image, 1, newwarp, 0.3, 0)

return result, left_fit, right_fit

```

0.4 Main Pipeline

0.4.1 Load Video and Extract frames

```

In [20]: # Load Video
clip = VideoFileClip("project_video.mp4")

# Get the frames from the video
frames = []
for x in clip.iter_frames():
    frames.append(x)

# Convert the sequence to numpy array
frames = np.array(frames)
print('Frame Sequence shape: ', frames.shape)

```

Frame Sequence shape: (1260, 720, 1280, 3)

0.4.2 Lane Finding Pipeline

```

In [21]: # Define Vertices:
vertices = np.array([(240,720), (580, 440), (700, 440), (1200,720)]), dtype

# To save
lane_map = []

print('Starting: ', end=" ")

# Get first frame
frame = frames[0, :, :, :]

# Correct for distortion
undist_frame = calibrate_image(frame, mtx, dist)

# Find the lane

```

```

result, left_fit, right_fit = first_lane(undist_frame, vertices)

# Append the final result
lane_map.append(result)

print ('>', end= " ")

# From Second frame to the end
for i in range(1, frames.shape[0]):
    # Get frame
    frame = frames[i, :, :, :]

    # Correct for distortion
    undist_frame = calibrate_image(frame, mtx, dist)

    # Find the lane
    result, left_fit, right_fit = second_to_end_lane(undist_frame, vertices)

    # Append the final result
    lane_map.append(result)

    # Progress Bar, every 100 frames.
    if i%100 == 0:
        print ('>', end= " ")

print ('end', end=" ")

```

Starting: > > > > > > > > > > > end

0.4.3 Save the output to output_images directory

```

In [22]: # Save the output image
print('Saving: >', end=" ")

for j in range(len(lane_map)):
    # Get the output image
    image = lane_map[j]

    # Filename to save
    name = 'output_images/'+(str(j))+'.jpg')

    # Save the image to directory
    imsave(name, image)

    # Increment counter
    j += 1

    # Progress Bar, every 100 frames

```

```

        if j%100 == 0:
            print ('>', end= " ")

    print ('end', end=" ")

```

Saving: > > > > > > > > > > > end

0.4.4 Convert Output Frames to Video

In [12]: # Define the codec and create VideoWriter object

```

import skvideo.io

writer = skvideo.io.FFmpegWriter("output.mp4", outputdict={'-vcodec': 'libx264'})

for i in range(1260):
    filename = 'output_images/' + str(i) + '.jpg'
    read = imread(filename, False, 'RGB')
    read = np.array(read)

    # Write to video
    writer.writeFrame(read)

writer.close()

```

In [3]: HTML("""
 <video width="320" height="200" controls>
 <source src="{0}">
 </video>
 """.format("output.mp4"))

Out[3]: <IPython.core.display.HTML object>